

Multi-Threaded Matrix Multiply

A Comparison Between Pthreads and OpenMP

Jesse Bannon

University of Washington, Tacoma

July 10, 2017

In this paper, I outline the performance characteristics between three implementations of dense matrix-matrix multiply: single-threaded, multi-threaded, and multi-threaded Strassen algorithm. Each implementation stores matrices row-major in memory and uses double-precision floating points. These metrics were gathered using an Intel i7 5820K CPU, which includes six cores and twelve threads using Intel's Hyper-Threading Technology. Additionally, in the multi-threaded implementations, we compare the difference in performance using Pthreads and OpenMP.

Single-Threaded One core performs a matrix-matrix multiply $C = AB$ using a trivial triple-nested for-loop.

Multi-Threaded Each thread t_i performs a subset of the matrix-matrix multiply $C_i = AB_i$, where B is partitioned by consecutive columns $B_i = \{b_j, b_{j+1}, \dots, b_{j+n}\}$ evenly amongst all t_i . In the case that B 's columns are greater than the number of threads, we assign a single column to a subset of threads and leave the remaining idle.

Strassen Algorithm Each operation (dense matrix-matrix add, subtract, multiply) of the Strassen algorithm is multi-threaded amongst all threads. Every matrix-matrix multiply recursively calls Strassen algorithm, inconsequently performing a depth-first traversal in the recursion. After the split input matrix size reaches a certain threshold, the multi-threaded implementation is used.

When multiplying a 1374-by-1374 matrix with itself using Pthreads, the 16-thread multi-threaded multiply performed best with a 2.46 second runtime. Each thread was responsible to multiply the matrix with its (approximately) 1375-by-85 subset. The performance comes from the memory lookups on B when multiply its columns. The Strassen algorithm, in this case, over-parallelized the problem size. Alternatively, using OpenMP, the 12-thread multi-threaded multiply performed best with a 2.70 second runtime. Pthreads slightly outperformed OpenMP in this case with a 10% speedup.

With Pthreads, the 3937-by-3937 matrix performs best using the 6-thread Strassen algorithm, with a 128-size stopping condition for dense matrix-matrix multiplies, in 45.11 seconds. The second best metric has the same configuration except it uses 12 threads. In this case, hyper-threading decreases performance. The problem size in this case is large enough to where the Strassen algorithm does not over-parallelize, which takes advantage of its lower complexity of $\mathcal{O}(n^{\log_2 7})$ compared to the traditional matrix-multiply complexity of $\mathcal{O}(n^3)$.

Surprisingly, OpenMP performed better than Pthreads. The 6-thread Strassen algorithm with a 64-size stopping condition took 38.21 seconds, which translates to a 20% speedup compared to using Pthreads.

Conclusion Different variants of multi-threaded dense matrix-multiply perform better depending on the problem size and underlying architecture. It's more efficient to use a traditional row-based matrix multiply on smaller sized problems compared to the Strassen algorithm, to mitigate over-parallelization overhead delay. Larger problem size's compute outweighs the cost of thread management, and can take




Figure 1: Elapsed time to perform dense matrix-matrix multiply using Pthreads

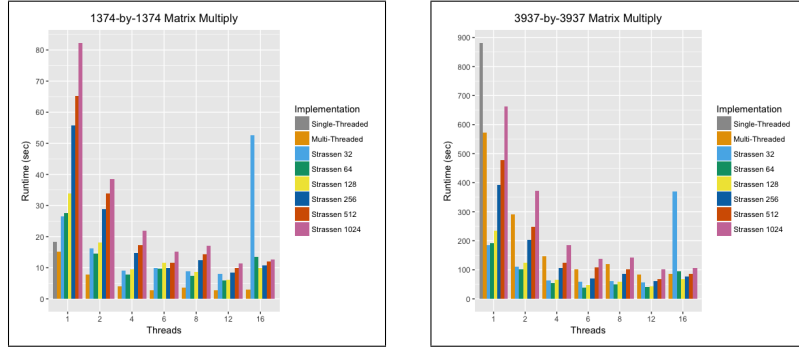


Figure 2: Elapsed time to perform dense matrix-matrix multiply using OpenMP

advantage the Strassen algorithm's lower complexity. Other factors such as cache line size and positioning of memory lookups between threads also affect performance, as seen in different sized dimension stopping conditions in the Strassen algorithm implementation.