

# Multi-Threaded Matrix Multiply

Jesse Bannon

University of Washington, Tacoma

July 2, 2017

In this paper, I outline the performance characteristics between three implementations of dense matrix-matrix multiply: single-threaded, multi-threaded, and multi-threaded Strassen algorithm. Each implementation stores matrices row-major in memory and uses double-precision floating points. These metrics were gathered using an Intel i7 5820K CPU, which includes six cores and twelve threads using Intel's Hyper-Threading Technology.

**Single-Threaded** One core performs a matrix-matrix multiply  $C = AB$  using a trivial triple-nested for-loop.

**Multi-Threaded** Each thread  $t_i$  performs a subset of the matrix-matrix multiply  $C_i = AB_i$ , where  $B$  is partitioned by consecutive columns  $B_i = \{b_j, b_{j+1}, \dots, b_{j+n}\}$  evenly amongst all  $t_i$ . In the case that  $B$ 's columns are greater than the number of threads, we assign a single column to a subset of threads and leave the remaining idle.

**Strassen Algorithm** Each operation (dense matrix-matrix add, subtract, multiply) of the Strassen algorithm is multi-threaded amongst all threads. Every matrix-matrix multiply recursively calls Strassen algorithm, inconsequently performing a depth-first traversal in the recursion. After the split input matrix size reaches a certain threshold, the multi-threaded implementation is used.




Figure 1: Elapsed time to perform dense matrix-matrix multiply

When multiplying a 1374-by-1374 matrix with itself, the 16-thread multi-threaded multiply performed the best with a 2.46 second runtime. Each thread was responsible to multiply the matrix with its (approximately) 1375-by-85 subset. The performance comes from the memory lookups on  $B$  when multiply its columns. The Strassen algorithm, in this case, over-parallelized the problem size.

The 3937-by-3937 matrix performs best using the 6-thread Strassen algorithm, with a 128-size stopping condition for dense matrix-matrix multiplies, in 45.11 seconds. The second best metric has the same configuration except it uses 12 threads. In this case, hyper-threading decreases performance. The problem size in this case is large enough to where the Strassen algorithm does not over-parallelize, which takes advantage of its lower complexity of  $\mathcal{O}(n^{\log_2 7})$  compared to the traditional matrix-multiply complexity of  $\mathcal{O}(n^3)$ .

**Challenges** The main challenge was implementing the Strassen algorithm. After a few hours of debugging, the implementation fell into place. Using the depth-first traversal allowed me to reuse my multithread matrix multiply which helped.

**Further Improvements** Parallelizing the Strassen algorithm in a way where multiple threads can work on different operations (multiply, add, subtract) simultaneously could perhaps yield better performance. This could potentially decrease memory-lookups by having multiple threads simultaneously fetch different parts of the matrices from main memory into the L3, where other threads would subsequently read that same data later.

**Locking** Dense matrix-matrix multiply does not require locking because the problem can be partitioned into completely independent sub-problems. As described above, each thread  $t_i$  is responsible for computing  $C_i = AB_i$ . Once all threads complete,  $C$  is formed. The closest thing to a lock in this problem is a barrier, to wait for all  $t_i$  to complete so we know  $C$  is formed.